Personal portfolio

Team 08 - Trung Nguyen – 2108488

Build role: Tester

Week 1

What I contributed

I looked up what qualifies as an SME, suggested the F&B industry because it seemed like a good choice, reminded the group that internal sub-question (for the firm we're writing the handbook for) should also be in the research proposal, and reviewed both drafts after they were made.

What my build role produced

I took on the tester role. There is nothing much this week yet for this role.

What I don't understand yet

I don't know how we're supposed to build AND operate an agentic CLI without any prior knowledge about programming.

Week 2

What the team published

We published a handbook page on geopolitics and vendor models for Giesen Roasting Solutions, a family-owned coffee roasting machine manufacturer. The page mentions that firms this size mostly buy their AI from other companies, so the real question is what happens if that company or its government changes course. It also explains what the EU Data Act gives a machine builder, including the right switch providers and take data out. It then ends with 5 first steps for an owner: a vendor list, three written questions to each vendor, an export test, scoring each vendor on dependency and price, and an exit clause in the next contract.

What I contributed

I helped double check the claims against their own source, except for the Dutch ones.

What the two gates asked, and what changed

What they asked:

- Reliability: Does having a source automatically mean a claim is accurate and relevant? How would you detect cases where the source is misleading or the AI misinterprets it?
- Data privacy: Your blueprint states that confidential interview material must never enter the Claude Project. Since this rule relies partly on team members following instructions and checking the project, what could happen if someone accidentally uploads confidential information?
- Memory: Your team uses session-log.md to remember previous sessions, with the project manager manually uploading updated logs. What problems might arise if the log is outdated, incomplete, or incorrectly updated? How could these problems affect the quality of the handbook?

What changed: I, as the tester, became more aware of the reliability of the sources, and I now check whether the original source supports the AI's interpretation. I also keep the protected folder with confidential information separate from the project.

What my build role produced

I revised the PRD and technical blueprint.

Week 3

What the team published

We made a page on Trust and AI safety. It argues that it is advised against trying to decide if "AI is safe" generally, and AI users should decide on their own: what data is allowed to go into AI tools, AI tools' actions within the system, and how open the staff is about their AI use.

What I contributed

As the tester, I helped double check the claims against their own source, except for the Dutch ones

What the two gates asked, and what changed

What the gates asked:

- What does your design do when the exact sentence is accurate but the claim built on it is not?
- How would you know the memory has failed before the 150-row trigger does?
- Data privacy: Which parts of "the line" are enforced by the system, and which depend on someone paying attention?

What I changed:

- How I sample: instead of opening 3 sources at random per chapter, I check for the ones that seem to carry the most weight in supporting the chapter's argument.
- With data: I treat confidential data with utmost care just as before.

What my build role produced

My build role didn't produce anything this week.

Week 4

What the team published

An unfinished page about SME knowledge gaps that's supposed to answer questions about SMEs' main barriers to AI and what the firm knows about its own delivery-date performance and planning data. It's incomplete because we still couldn't book a visit for an interview.

What I contributed

I offered to create an operations manager persona to interview in case our request for a real one gets denied/no responses, since the idea was green lit by one of the lecturers Witek.

What was asked, and what changed

I didn't attend it because I wasn't feeling well, but from my groupmates, the questions asked by the lecturers were about our tool, page, and interview guide: a demo for the tool, what makes the tool agentic and which element had the greatest impact, whether anything didn't work with the tool, and field research progress.

What my build role produced

My build role didn't produce anything this week.

Week 5

What the team published

A completed page on SME knowledge gaps: SMEs are adopting AI (45% in 2025, up from 20% in 2023); Lack of experience, among privacy and unclear legal consequences are what's stopping firms from adopting AI; 83% of Dutch SME manufacturers still aren't aware of how and which process suit AI.

A page on AI use cases at SMEs: Most firms use AI in the office, mainly for marketing, sales, and admin, only 18% of Dutch AI-using firms use AI in production. A study of 2 German make-to-order firms suggests that predicting the delivery date at the quote can lower the risk of late delivery, but it carried 35% more error than the planners' estimation.

A demo on how to run our Agentic CLI.

What I contributed

I helped double check the claims against their own source, except for the Dutch ones for the SME knowledge gaps page. I couldn't do it for the other page due to illness.

Since the real interview never happened due to no response from Giesen, I asked the team if we could create an AI persona and put the same questions to it, and we all agreed. I later set up an AI persona and interviewed it.

I helped with editing the demo.

What the two gates asked, and what changed

What was asked:

- Your PRD measures how many sentences trace to an accepted source. Does that measure what makes this page good?
- What would have to be true for your pick to be wrong, and where does the page say so? The page says it found no documented, measured case of AI planning at a Dutch SME machine builder. The one comparable study is two German firms, not food-processing machinery. In that study the early estimate had 35% more error than the planners', and some errors no one could correct.

What changed:

- I consider using another measure alongside the 95%, for example the share of recommendation sentences that are labeled as interpretation.

- For me as the tester, this means that I need to check the interpretation steps as well as the sources.

What my build role produced

One interview transcript with an AI persona and a complete demo video.