

# Product Requirements Document (First Draft)

*AEL — Agentic platform for the AI handbook*

## 1. Problem and Users

**Problem:** Producing a rigorously fact-checked, localized AI handbook in six weeks requires evaluating a high volume of regional data (KvK, CBS, Dutch trade press and vendor documentation). Manually retrieving, translating, and synthesizing this data limits how deep our research can go.

**Users:** The direct users are our four-person student research team. The downstream beneficiary is the SME reader, Koffiebranderij Peeze, or a comparably sized equipment-service SME, who relies on the handbook being accurate and actually usable.

## 2. Functional Requirements (What it must let a user do)

The agentic platform must function as an autonomous research assistant that lets a user complete the path from raw source to a published handbook page:

1. Prompt: input a weekly research theme (e.g. "SME predictive maintenance vendors Europe").
2. Retrieve & Cross-Reference: command the agent to search the web, retrieve relevant Dutch and regional market data, and cross-reference vendor claims against relevant EU regulation (e.g. the AI Act) where it applies.
3. Evaluate: automatically appraise retrieved sources against our Handbook Quality Criteria, discarding anything enterprise-focused, excessively costly, or clearly non-compliant.
4. Draft: synthesize the surviving material into a structured Markdown chapter with inline, verifiable citations, ready to publish to our wiki site.

## 3. Output Qualities

A successful handbook page generated by the platform must be:

- Traceable: every factual claim, price estimate, or regulatory mention has a verified original source
- Filtered for scope: strips out AI “hype” and stays focused on the physical and financial reality of a company Peeze's size, roughly 50 to 60 staff
- Editorially clean: drafted in professional English, following the course's chapter anatomy template, using markdown tables and bolding for scannability

## 4. Known Constraints & The Language Gap

- Agentic CLI: the tool is built by directing an agentic command-line interface, not by hand-coding Python scripts
- Accessibility: a non-coder on the team (e.g. the Product Manager) must be able to run the pipeline via clean terminal commands
- Manual fallback: every automated step (search, translate, draft) outputs a readable log so we can manually complete it if the agent breaks or hits a rate limit
- The Dutch language gap: critical context, Peeze interview transcripts, Dutch regional press, CBS and KvK data, and whatever accounting or ERP documentation turns out to be relevant once we've confirmed what Peeze actually runs, will be in Dutch. The platform must include an automated translation layer: the agent translates Dutch sources to English before evaluation, and saves the raw translation to a text file so non-Dutch-speaking teammates can audit its logic

## 5. Out of Scope (For Now)

- Building, training, or deploying an actual predictive maintenance algorithm for Peeze
- Automated, direct-to-web publishing (human review and the Socratic/peer gates happen before anything is published)

- Audio processing (field interview recordings are transcribed externally before text ingestion)

## 6. Open Questions and Assumptions

- Assumption: the LLM powering our agent has a large enough context window to digest long, dense documents (like EU AI Act summaries) without losing detail
- Open question (financial hallucination): B2B vendors often hide pricing. How do we build an evaluation check that stops the agent from inventing SaaS costs when it can't find a public price?
- Open question (API limits): what are the latency and token-limit realities of chaining search, translation, and drafting in a single run?
- Open question (decision-maker): now that Peeze is part of Coffee Growth Group, is a tool like this Peeze's decision to make alone, and does that change who the handbook's “downstream beneficiary” really is?